

Distribution-Robust Vector Search: Calibrated Routing, Coverage, and Precision-Adaptive Scanning for IVF-PQ

Thomas Dybdahl Ahle
Normal Computing
thomas@normalcomputing.com

Abstract

Scalable approximate nearest-neighbor (ANN) search routes a query to a few partition cells, scans quantized codes inside them, and re-ranks a shortlist. On *out-of-distribution* (OOD) workloads — where queries come from a different distribution than the database, as in cross-modal (text→image) retrieval under inner product — we first *localize the bottleneck*: a battery of scan-side and code-side improvements (wider SIMD, higher-resolution and rotated codes, RaBitQ, ADC routing) leave OOD recall/QPS essentially unmoved, while the achievable recall is capped by which cells get *routed to and covered*. We then present a collection of routing- and coverage-directed techniques, each measured in isolation: a one-parameter routing recalibration (γ) grounded in a cap-versus-ball selectivity argument; k -NN-graph coverage augmentation; a joint cross-level partition optimization; multi-assignment; and a corollary that *reverses* standard practice — once graph coverage is present, coarser cells win. Composed and evaluated under a same-hardware, tight-pairwise protocol against ScaNN at its official configuration, the engine attains higher QPS at matched recall@10 at 1M and 10M, extends to 100M — where no ScaNN configuration builds within a 2h gate — and transfers to the streaming track. We then turn to the *in-distribution* regime (Cohere-768, cosine), where the same engine initially loses to HNSW by $2\times$ — and localize *that* bottleneck too: not the architecture but a *scan-precision policy*. The int8 saturating fast-scan caps each subspace lookup table at ~ 4 bits; its quantization noise grows as $\sqrt{m}$ with the subspace count, harmless at $m=100$ ($d=200$) but fatal to the small cosine margins at $m=384$ ($d=768$), where in-pool ordering collapses to 0.07. Switching the same codes to an int16 scan restores ordering to the pool ceiling (0.87), and precision auto-selection by m — plus the coverage levers above — makes the engine win in HNSW’s home regime: it beats HNSW at 1M *graph-free*, and at 10M/ $d=768$ (HNSW’s strongest setting) its query path beats HNSW across the recall range (~ 1.2 – $1.9\times$, reaching recall HNSW cannot) after a minutes-scale internal k -NN-graph build — the engine *self-builds* the graph by NN-descent (its routed self-search is coverage-capped at $d=768$ and cannot; local join sidesteps routing entirely), so the win is fully self-contained. We report the negative results as first-class evidence for both localizations, and treat the measurement methodology — which repeatedly caught mirages during development — as a contribution.

1 Introduction

Approximate nearest-neighbor (ANN) search is the retrieval primitive under retrieval-augmented generation, semantic search, and recommendation. At scale the dominant family is inverted-file with quantization (IVF-PQ): partition the database into C cells, at query time *route* to the $p \ll C$ most promising cells, *scan* compressed codes within them, and exactly *re-rank* a shortlist. A large literature has driven each stage — anisotropic quantization and SIMD fast-scan for the scan, product and residual codes for memory, graph indices as an alternative to IVF — overwhelmingly evaluated where queries and database are drawn from the *same* distribution.

We study *distribution robustness*: one engine, measured in two regimes that break different stages of the pipeline. The first is *out-of-distribution* (OOD) search, where the query distribution differs from the database. This is not a corner case: cross-modal retrieval embeds a text query and searches image vectors

under inner product, and the big-ANN [9] “text2image” track (image base, text queries, $d=200$) is exactly this setting. It is where the strongest systems separate, and — we argue — where the conventional intuition about what to optimize is wrong. The second is the *in-distribution, high-dimensional* regime (Cohere-768 embeddings under cosine), the bread-and-butter workload of graph indices like HNSW [3] — where the same IVF-PQ engine initially loses by $2\times$, for a reason that turns out to be neither routing nor architecture but arithmetic precision in the scan (§11).

Our organizing claim is a *bottleneck localization*: on OOD workloads the binding constraint is *routing coverage* — whether the true neighbor’s cell is among the p probed and its members reachable — not scan throughput or code fidelity. We support this from both sides. A battery of scan- and code-side improvements (wider SIMD, higher-resolution and rotated codes, RaBitQ [8], ADC routing [10]) leaves OOD recall and QPS essentially unmoved (§13); meanwhile a small set of routing- and coverage-directed techniques each yields a measured gain (§6–§10). We therefore present the work as a measured collection — every technique reported with the effect it produces in isolation, and the negatives kept as first-class evidence for the localization. Composed and evaluated under a same-hardware, tight-pairwise protocol against ScaNN [1] at its best configuration, the result attains higher QPS at matched recall@10 at 1M and 10M (and, where ScaNN cannot build within a 2h gate, remains the only method to index 100M), and transfers to the streaming track.

Contributions (each measured in isolation).

- **A bottleneck localization** (§5): scan- and code-side levers do not move OOD; routing/coverage does. The negatives (§13) are the evidence.
- **γ routing recalibration** (§6), grounded in a cap-vs-ball selectivity argument (§4). *Effect: $\approx$ half the probes at matched OOD recall, zero query-time cost.*
- **k -NN-graph coverage** (§7). *Effect: composes with γ ; at 1M moves the ScaNN ratio $1.18 \rightarrow 0.98$.*
- **The coarse-cell corollary** (§8). *Effect: with graph coverage present, coarser cells win; at 10M $1.01 \rightarrow 0.85$.*
- **Joint cross-level partition optimization (tree-EM)** (§9). *Effect: $+0.6$ – 1.0 recall@10 points at 10M, QPS-neutral.*
- **Multi-assignment (SOAR [2])** (§10). *Effect: raises the routing recall ceiling to baseline parity.*
- **Scan precision as a function of code length** (§11): the int8-saturating fast-scan’s per-subspace ~ 4 -bit LUT cap injects noise growing as $\sqrt{m}$; at $m=384$ (in-distribution $d=768$ cosine) it collapses in-pool ordering to 0.07 while the *same codes* under an int16 scan rank to the pool ceiling (0.87). *Effect: with precision auto-selected by m (plus an AVX-512 int16 kernel and graph coverage), the engine goes from $2\times$ behind HNSW on Cohere-768 to a graph-free win at 1M and, after a minutes-scale internal NN-descent graph build, a ~ 1.2 – $1.9\times$ win at 10M — without abandoning the IVF-PQ cascade.*
- **Systems levers for streaming** (§12): batch-parallel insert ($4.5\times$), low-live adaptive p (*worst step* $0.74 \rightarrow 0.94$), and a parallelization fix ($7.6\times$ at 8 threads).
- **A same-hardware evaluation methodology** (§14) and an honest scope of what the ratios do and do not claim.

2 Background: Standard Building Blocks

Our system is assembled from well-established components; we summarize them here so the contributions in §3–§10 can be stated as deltas against known practice.

Inverted file (IVF). Partition the database into C cells (the Voronoi regions of C centroids). At query time, *route* the query to its $p \ll C$ nearest centroids and scan only those cells. End-to-end recall is capped by whether the true neighbor’s cell is among the p probed — the routing question that this paper argues is the OOD bottleneck.




Figure 1: Query pipeline (top) and the offline-built index it consumes (bottom, dashed). Edge labels are the shrinking candidate set: p probed leaves $\rightarrow$ top- t pool $\rightarrow K$ cascade survivors $\rightarrow$ top-10. γ calibration, the k -NN graph union, and the coarse-cell tree geometry are this work’s contributions (Table 1).

Hierarchical routing. Scoring all C centroids per query is $O(Cd)$, which dominates at scale (fine C is needed to keep cells small). A *tree* of centroids — coarse k -means over C_0 centroids, then fine k -means within each — reduces routing to $O(LC^{1/L})$ for an L -level tree (FAISS [4] uses an HNSW graph over the centroids for the same purpose). We use hierarchical k -means.

Product quantization (PQ). Split each vector’s d dimensions into m subspaces and quantize each subspace to a small codebook (16 codewords for 4-bit PQ). Approximate distances are computed by lookup-and-add over per-subspace tables (ADC). With 4-bit codes the table lookup is a register shuffle (`vpshufb`), enabling the SIMD “fast-scan” kernels (Quick-ADC, ScaNN, FAISS FastScan). *Anisotropic* / score-aware quantization (ScaNN) weights the quantization loss to preserve inner-product ranking rather than reconstruction.

Exact re-ranking. The quantized scan is used only to select a shortlist; exact distances (int8 or float, read for the few shortlisted vectors) re-order it. The shortlist depth t trades recall for cost.

Multi-assignment. Replicating each point into its top- a_0 cells (a space/time tradeoff) raises the chance a query’s neighbor lands in a probed cell.

The OOD workload. The big-ANN OOD track (text2image) pairs *image*-embedding base vectors ($d=200$) with *text*-embedding queries under inner product, with a float ground truth. Base norms are concentrated but not constant (≈ 0.81 – 0.99); queries are unit-norm. This distribution gap between base and queries is what breaks L_2 -trained routing (§6).

3 System Overview

Figure 1 shows the pipeline; we describe the build and query paths, then delineate standard vs. contributed components (Table 1).

Build path.

1. **Quantize.** Float base $\rightarrow$ int8 via a single global scale ($127/\max|x|$); the int8 base is the native scan/route representation.
2. **Train the routing tree.** Hierarchical k -means: C_0 coarse centroids, then fine centroids within each coarse cell to K_f leaves total (2-level; generalizes to L levels).
3. **Joint refinement (tree-EM, §9).** Optionally re-optimize all levels jointly (E: reassign sample points to their nearest leaf via the beam descent; M: recompute every level’s centroids).
4. **Multi-assign (SOAR, §10).** Assign each point to a_0 leaves with spill that covers the residual direction of the first.

Algorithm 1: Build

Input: float base $X_f \in \mathbb{R}^{n \times d}$; tree sizes C_0, K_f ; multi-assign a_0 ; calibration γ ; EM rounds R ; graph degree k

Output: index $(\{Q_0\}, \{c_j\}, b, \text{PQ blocks, raw, } G)$

```
1  $\sigma \leftarrow 127 / \max_i \|X_f[i]\|_\infty$ ;  $X \leftarrow \text{int8}(\sigma X_f)$ ; // global-scale quantize
2  $\{Q_0\} \leftarrow k\text{-means}(X, C_0)$ ;  $\{c_j\} \leftarrow \text{fine } k\text{-means within each coarse cell to } K_f \text{ leaves}$ ;
3 for  $R$  rounds // tree-EM (§9); optional do
4   E: assign each sample to its nearest leaf via  $\text{Route}(\cdot, 1)$ ;
5   M: recompute  $\{Q_0\}, \{c_j\}$  from the assignments;
6 foreach point  $x_i$  // SOAR multi-assign (§10) do
7   assign  $x_i$  to its top- $a_0$  leaves; spill covers the residual of the 1st;
8 per leaf: encode members as 4-bit anisotropic PQ, pack into 16-row vpshufb blocks; store raw int8
   in leaf order;
9  $b_j \leftarrow (\gamma - 1)\|c_j\|^2$  for all  $j$ ; // routing bias
10  $G \leftarrow \text{base-side } k\text{-NN graph (self-search; §7)}$ ;
11 return index  $(\{Q_0\}, \{c_j\}, b, \text{PQ blocks, raw, } G)$ ;
```

Algorithm 2: Route (tree descent) – returns the p best leaves

```
1 Procedure  $\text{Route}(q, p)$ :
2    $S_0 \leftarrow \text{top-}b_0 \text{ coarse cells minimizing } s_\gamma(q, Q_0)$ ; // int8/VNNI kernel
3    $\mathcal{L} \leftarrow \text{fine children of } S_0$ ;
4   score each  $c_j \in \mathcal{L}$  by  $\|c_j\|^2 - 2q \cdot c_j + b_j$ ;
5   return the  $p$  leaves of  $\mathcal{L}$  with smallest score;
```

5. **Encode.** 4-bit anisotropic PQ ($m=d/2$ subspaces); pack codes into cell-contiguous 16-row blocks laid out for the **vpshufb** fast-scan. Store the raw int8 vectors in cell order too (cache-warm re-rank gathers).
6. **Calibrate & augment.** Precompute the per-cell γ routing bias (§6); build the base-side k -NN graph sidecar (§7).

Query path.

1. **Route (tree descent).** Score the C_0 coarse centroids (int8, VNNI kernel), keep the top- b_0 ; expand to their fine children, score those with the γ bias added, and return the top- p leaves.
2. **Scan.** Fast-scan the 4-bit PQ codes of the p probed cells (**vpshufb** ADC with int16 accumulation for ~ 12 -bit selection resolution; optional AVX-512 64-wide), collecting a bounded top- t pool.
3. **Cascade re-rank.** Deduplicate the pool by original id (needed under multi-assignment); exact-rescore the t survivors in int8 (VNNI dot); prune to the top- K ($K=16$).
4. **Float re-rank + graph union.** Compute exact float inner products for the K survivors (reading the original float vectors) unioned with the k -NN graph neighbors of the top- M ; return the top-10.

Algorithms 1–3 state the pipeline precisely. All routing scores use the calibrated form $s_\gamma(q, c) = \gamma\|c\|^2 - 2q \cdot c$ (§6); the constant $\|q\|^2$ is dropped, and smaller s_γ means a better cell. Leaf centroids are $\{c_j\}_{j=1}^{K_f}$ with precomputed bias $b_j = (\gamma - 1)\|c_j\|^2$ so that $s_\gamma(q, c_j) = \|c_j\|^2 - 2q \cdot c_j + b_j$ is one dot product plus an additive constant.

Algorithm 3: Query

Input: int8 query q , float query q_f ; probes p ; pool depth t ; prune K ; graph fan-in M

```

1  $\mathcal{C} \leftarrow \text{Route}(q, p)$ ;
2  $\mathcal{P} \leftarrow \emptyset$ ; // bounded top- $t$  pool
3 foreach  $cell \in \mathcal{C}$  do
4   fast-scan its 4-bit PQ codes (vpshufb ADC, int16 accum)  $\rightarrow$  approx. scores;
5   keep the  $t$  best in  $\mathcal{P}$  (threshold-bounded collect);
6  $\mathcal{P} \leftarrow$  dedup  $\mathcal{P}$  by original id ; // needed under  $a_0 > 1$ 
7  $\mathcal{K} \leftarrow$  top- $K$  of  $\mathcal{P}$  by exact int8 IP (VNNI dot) ; // cascade prune
8  $\mathcal{R} \leftarrow \mathcal{K} \cup \bigcup_{r \in \text{top-}M(\mathcal{K})} G[r]$  ; //  $k$ -NN-graph coverage union
9 return top-10 of  $\mathcal{R}$  by exact float IP  $\langle q_f, x_r \rangle$ ;
```

Table 1: Components of the pipeline: standard building blocks vs. this work.

Component	Basis	Our delta
int8 IVF partition	standard	—
Hierarchical k -means tree	standard (FAISS-style)	fast int8/VNNI descent kernel
4-bit anisotropic PQ	ScaNN / FastScan family	int16-resolution + 32/64-wide fast-scan
Multi-assignment (SOAR)	ScaNN-derived	dedup-before-cap correctness fix
int8 $\rightarrow$ float re-rank cascade	folklore	the specific $K=16$ composition
γ routing calibration	ours	the OOD miscalibration fix
k -NN-graph coverage union	ours	composition with γ + float re-rank
Coarse-cell geometry knee	ours	reverses “finer is better”
Tree-EM joint optimization	ours	cross-level partition objective
Batch-insert / adaptive- p (streaming)	ours	§12

4 Theoretical Grounding (Scoped)

Under concentration, additive ball filters become non-selective — in high dimension almost everything falls within $R+r$ — whereas angular (spherical-cap) filters remain selective, and k -means on centered data realizes exactly such a data-dependent angular partition. This framing motivates γ : the routing score $\gamma\|c\|^2 - 2q \cdot c$ interpolates from an additive ball ($\gamma=1$, L_2) to a pure inner-product cap ($\gamma \rightarrow 0$), so the geometry predicts that $\gamma < 1$ is needed precisely in the concentrated OOD regime — which our experiments confirm. We build on two principles from this line of work: *route-precise*, *rank-cheap*, and *multi-assignment as a space/time tradeoff*. We do *not* claim the recursive-tree polylog-loss result, which does not pay off at the $\leq 10\text{M}$ scales evaluated here.

5 Localizing the OOD Bottleneck

We separate two ceilings. The *pool-recall ceiling* is the recall obtainable by exactly re-ranking the *entire* probed pool — it depends only on routing (which cells, hence which points, are reachable), not on the scan’s approximation. The *scan-limited recall* is what the actual (quantized) scan plus shortlist delivers. If the two coincide, the scan is already extracting everything the routed pool contains and there is no scan headroom; recall is then a pure function of routing coverage.

On OOD text2image this is what we observe: at a fixed probe budget the scan-limited recall sits at the pool-recall ceiling, and that ceiling — not the scan — is what moves when we change routing. Two consequences follow, and the rest of the paper is their systematic exploration. First, techniques that improve *coverage* (route calibration, graph augmentation, joint partition optimization, multi-assignment) should raise recall; §6–§10 confirm each does, with a measured effect. Second, techniques that only improve the *scan or the code* (wider SIMD, higher-resolution or rotated codes, alternative quantizers, ADC routing) should not help, because there is no scan headroom to recover; §13 confirms each does not. The negatives are thus not

incidental — they are the control that localizes the bottleneck to routing coverage.

6 γ : One-Parameter Routing Recalibration

An IVF router ranks cells by a score over the query q and centroid c . For L_2 k -means the score is $\|q - c\|^2 = \|q\|^2 + \|c\|^2 - 2q \cdot c$; dropping the query-constant term leaves $\|c\|^2 - 2q \cdot c$. We write the family

$$s_\gamma(q, c) = \gamma\|c\|^2 - 2q \cdot c, \quad (1)$$

where $\gamma=1$ recovers L_2 routing and $\gamma \rightarrow 0$ recovers pure inner-product (angular) routing. Under distribution shift these disagree in a specific way: the $\|c\|^2$ penalty demotes high-norm centroids, but for OOD queries the inner-product-relevant neighbors concentrate precisely near those high-norm cells, so $\gamma=1$ systematically misroutes them. A single global $\gamma < 1$ rebalances the two terms. It is grounded in the cap-vs-ball argument of §4: γ slides routing along the additive-ball ($\gamma=1$, non-selective under concentration) to spherical-cap ($\gamma \rightarrow 0$, selective) axis.

Crucially the fix is *free at query time*: s_γ differs from the L_2 score only by the per-cell constant $(\gamma-1)\|c\|^2$, which we fold into a precomputed integer bias added during cell scoring (no extra query work, no index-format change). We fit γ once per dataset; $\gamma \approx 0.5$ is optimal on text2image at both 1M and 10M, and the recall is flat across $\gamma \in \{0.4, \dots, 0.6\}$ near the optimum.

Effect. Reaches matched OOD recall at $\approx$ half the probes (p : 54 $\rightarrow$ 29 at 1M; $\approx 2.5\times$ probe cut at 10M). Zero query-time cost.

7 k -NN-Graph Coverage Augmentation

Routing to p cells can still miss a true neighbor whose cell went unprobed. We recover such misses with a graph, without probing more cells. Offline we build a base-side inner-product k -NN graph (a per-point adjacency list; constructed by ScaNN self-search, by the engine querying itself, or — where routed self-search is coverage-capped — by the engine’s internal NN-descent pass, below; the first two agree on $\sim 92\%$ of edges). The graph is a method-agnostic offline artifact, so the query-time comparison is unaffected by which builder is used, and we verified this empirically: on OOD text2image-10M at the champion config, swapping the ScaNN-built graph for the engine’s own self-built one changes recall by $\leq 0.1\text{pt}$ and QPS within noise. Where the source could matter for a claim we use the self-built graph outright — e.g. the 35M in-distribution result (Table 8) uses the engine’s *own* graph and still wins (there it trails a competitor-built graph by $\sim 0.4\text{pt}$ but keeps the win), so no comparison borrows anything from the opponent. At query time, after the scan yields a shortlist, we take the top- M shortlist members and *union their graph neighbors into the exact re-rank set*. The intuition is transitive: a shortlisted point near the query is itself near the true neighbor, so the true neighbor tends to appear among that point’s stored nearest neighbors even when its own cell was not probed. This adds coverage that is *independent of the partition* — the property that drives the coarse-cell corollary (§8). It composes with γ : the two attack different misses (calibration fixes which cells are ranked highly; the graph reaches beyond the probed cells), overlapping only partially. The union is deduplicated against the scanned pool (already-scanned neighbors are not re-scored) and the adjacency reads are software-prefetched, so the added cost is a small, bounded re-rank increment.

Effect. At 1M, composed with γ , the ScaNN ratio drops $1.18 \rightarrow 0.98$ (the first sub-1 result); at 10M it lets t_{surv} halve at matched recall. Graph parameters scale with n : fan-in M and pool depth grow from $M \approx 25$ at 1M to $M=64$ at 100M as distractors multiply.

Repair depth (and where it helps). The union above is one hop. Iterating it — expand the current best- M members’ neighbors, re-score, reselect, for R rounds — is a *batched* graph walk (beam width M ; the whole cohort is SIMD-rescored and a single exact re-rank closes it, so no per-query priority queue). It has a clean cost/quality model: each hop closes a geometric fraction ($\approx 56\%$) of the still-repairable miss mass, while p and t_{surv} raise a per-probe reachability *ceiling* that hops cannot exceed (fit to $R^2=0.99$ on a 20-point grid, so the two knobs tune independently, no black-box search). Crucially, extra hops are a matched-recall win *only above a ≈ 0.91 – 0.92 crossover*: at the leaderboard’s recall@10 ≥ 0.90 gate the single-hop point is $\approx 6\%$

faster (a second hop reaches 0.90 only by *overshooting* recall at a QPS cost), whereas for high-recall targets (0.94–0.96) a second or third hop reaches them at far lower p than one hop. We therefore keep hop count a knob (default 1), not a default — a high-recall lever, not a gate-level one. A best-first frontier (expand the globally-best unexpanded node) adds +0.3–0.8 pt *in-distribution* but is neutral OOD: base-metric edges do not align with off-manifold queries, so smarter reselection finds no better seeds. Reporting this required comparing at the recall *gate*, not fixed p — a fixed- p “+1.7 pt” reads as a win but is not one at the gate.

Self-building the graph: internal NN-descent. How the graph is *built* matters for self-containment. The engine’s *routed* self-search cannot build a good graph at 10M/ $d=768$: routing is coverage-capped (§17), and no practical probe depth fixes it (0.63 overlap@16 even probing 6% of cells). NN-descent [7] sidesteps routing entirely: starting from a random (or cheaply routed) adjacency, each round joins every point against its neighbors-of-neighbors — forward and (capped) reverse edges — and keeps the top- k ; convergence relies on transitivity of proximity, not on any partition. Our implementation is a race-free *pull* variant (each point updates only its own row; int8 VNNI dots; per-thread stamped-hash dedup) with two fixes that matter at inner-product hubness: the capped reverse lists are *reservoir-resampled each round* (a first-come cap freezes hub in-neighbors and permanently misses their pairings), and inserted edges stay “new” for two rounds (one round of gating loses join chances at saturated hubs).

Effect. Cohere 10M ($d=768$): random init reaches 0.864 overlap@16 in 24 min (16 threads); seeding with the engine’s own cheap routed self- k NN reaches **0.904** in 8 min of descent — query-time equal-or-better than the multi-hour external near-exact graph (recall 0.9871 vs 0.9869 at matched probes; 0.9904 vs 0.9901 at the top point). At 1M: 0.895 in ~1 min. The 10M win is therefore fully self-contained.

8 Routing Cost and the Coarse-Cell Corollary

8.1 A routing cost model

Consider an L -level tree over K_f leaves with cell counts $C_0 < C_1 < \dots < C_{L-1} < C_L = K_f$ and per-level beam widths $b_0, \dots, b_{L-2}$ (the number of cells expanded when descending from level i to level $i+1$). Routing a query costs, in centroid distance evaluations,

$$\text{route}(q) = C_0 + \sum_{i=1}^{L-1} b_{i-1} \frac{C_i}{C_{i-1}}, \quad (2)$$

since the coarse level scores all C_0 centroids and level i expands the b_{i-1} best parents from level $i-1$, each contributing C_i/C_{i-1} children. The scan then reads $\text{scan}(q) = p \cdot (n/K_f) \cdot a_0$ candidates (p probed leaves $\times$ points/leaf $\times$ multi-assignment). For a 2-level tree (2) is $C_0 + b_0 K_f / C_0$, minimized at $C_0 = \sqrt{b_0 K_f}$ to give $2\sqrt{b_0 K_f} = O(\sqrt{K_f})$; with L levels and balanced branching this falls to $O(L K_f^{1/L})$. The two costs pull oppositely in K_f : finer leaves shrink the scan (n/K_f) but grow the routing. The beams b_i are the “routing aggression” knob — wider beams retain the true cell with higher probability at linear cost — and OOD queries, which route poorly, need wider beams (§15).

Table 2 gives geometry and routing cost by scale. Two levels suffice through 10M; at 100M a 2-level tree at its optimal $C_0 = \sqrt{b_0 K_f} \approx 16,400$ would cost $\approx 33,000$ evals, whereas the 3-level tree roughly halves this (14,336) — so the third level is introduced only at 100M, where routing is a large enough share of the query to matter.

8.2 Scaling laws for geometry

To set geometry at scale without per- n retuning, we sweep K_f at $n \in \{10^5, 3 \cdot 10^5, 10^6, 3 \cdot 10^6, 10^7\}$ (2-level, graph off), at each n finding the minimum p that reaches $\text{recall@10} \geq 0.90$ (deterministic, hence load-independent) and the resulting cost route + scan from Eq. (2). The cost-minimizing geometry follows clean power laws (log-log fit):

$$K_f^* \sim n^{0.66}, \quad \frac{n}{K_f^*} (\text{pts/leaf}) \sim n^{0.34}, \quad p^* \sim n^{0.18}, \quad C_0^* \sim n^{0.33}.$$

Table 2: Tree geometry and per-query routing cost (Eq. 2) by scale. Beams are aggressive ($b_0 \approx 12\text{--}25\%$ of coarse cells) because OOD queries route poorly. The 100M geometry is validated by the routing-cost scaling law (§8.1).

scale	L	K_f	cell counts	beams	route evals	pts/leaf
1M	2	16,384	$C_0=768$	$b_0=96$	$768 + 96 \cdot 21 = 2,816$	61
10M	2	65,536	$C_0=4096$	$b_0=128$	$4096 + 128 \cdot 16 = 6,144$	152
100M	3	524,288	$C_0=2048, C_1=32768$	$b_0=512, b_1=256$	$2048 + 512 \cdot 16 + 256 \cdot 16 = 14,336$	190

The central law is $\text{pts/leaf} \sim n^{1/3}$: optimal leaves *grow* with n , because routing $O(\sqrt{K_f})$ rises while the scan $O(pn/K_f)$ is held down by the growing K_f — the coarse-cell corollary expressed as a scaling law. Probes grow only as $n^{1/5}$. Extrapolating, $n=10^9$ calls for $K_f \approx 10^6$ (~ 1000 pts/leaf). These calibrate the 2-level baseline; graph coverage (§7) shifts the optimum coarser and a third routing level (cheaper, $O(K_f^{1/3})$) shifts it finer, so the 100M geometry is validated by direct A/B (§15) rather than read off the 2-level law.

8.3 The corollary

The cost model (2) explains *why*. Without graph coverage, recall demands fine K_f (small leaves make the true neighbor’s leaf likelier to be among the p probed), paying $O(\sqrt{K_f})$ routing. The k -NN-graph union (§7) supplies that coverage *independently* of K_f , so the tree can use *coarser* leaves — cheaper routing — at fixed recall. The optimum is a knee: too fine wastes routing, too coarse inflates the scan ($p \cdot n/K_f$). Empirically the knee sits near 150 points/leaf at 10M, and both finer and coarser builds lose to it.

Effect. At 10M, moving from the fine-cell to the 152-point-cell geometry takes the ScaNN ratio $1.01 \rightarrow 0.85$; routing evals/query drop $\sim 12,000 \rightarrow 6,144$.

9 Joint Cross-Level Partition Optimization (Tree-EM)

Hierarchical k -means trains each level greedily top-down, so a leaf centroid is optimized for its parent’s Voronoi cell, not for the query-time beam descent that actually reaches it. We instead optimize all levels jointly by EM against the search objective: the E-step assigns each sample point to the leaf its *beam descent* lands in (not its globally-nearest leaf), and the M-step recomputes every level’s centroids from those assignments. Iterating co-adapts the levels so the descent lands better. The E-step is itself an ANN query against the tree’s own centroids, so it runs at a shallow beam (a few probes, not the full query beam), making EM cheap; it is a build-time cost only and adds nothing at query time. It composes with γ and the graph (it improves the partition; they improve routing and coverage on top).

Effect. $+0.6\text{--}1.0$ recall@10 points at 10M at matched probe, QPS-neutral (build-time only).

10 Multi-Assignment (SOAR)

The remaining coverage lever is redundancy: assign each point to its top- a_0 cells rather than one, so a query reaching any of them finds the point. Naïve replication puts near-duplicate copies in adjacent cells; SOAR instead chooses the extra cells to cover the *residual direction* of the first assignment, so the copies span different query directions and buy more coverage per unit storage. Realizing the gain required a correctness fix: the exact-re-rank shortlist must be deduplicated by original id *before* the survivor cap, otherwise the a_0 duplicate copies of a point crowd the fixed-size shortlist and starve distinct ids (recall then falls, and paradoxically decreases with more probes — the signature of the bug). With the fix, multi-assignment raises the routing-coverage ceiling.

Effect. $a_0=2$ reaches OOD coverage parity with the baseline partition; higher a_0 extends the high-recall frontier at a storage/scan cost (quantified in table).

11 Scan Precision Must Scale with Code Length

The techniques so far repair *routing and coverage*, the OOD bottleneck. The in-distribution, high- d regime breaks a different stage, and localizing it is a second instance of the paper’s method: measure stages in isolation until the failing one confesses.

The symptom. On Cohere-768 (10M, cosine, unit-norm) the engine needed an exact-re-rank pool of $t \approx 32,000$ to reach recall 0.90 — $60\times$ the OOD operating depth — and recall *decreased* with more probes. Stage-isolation pinned it: the routed pool *contains* the true neighbor (pool ceiling 0.87–0.94), and an 8-bit re-ranker recovers exactly the ceiling, but the 4-bit fast-scan’s *own ordering* of the pool is catastrophic — the true neighbor almost never ranks in its top-100 (in-pool ordering recall 0.07).

The cause is arithmetic, not codes. The SIMD fast-scan accumulates lookup-table (LUT) values in *saturating int8*; to avoid saturation within a hoist group, each subspace’s LUT is quantized to a ~ 4 -bit range (cap $\lfloor 120/\text{HOIST} \rfloor = 15$). The resulting per-subspace quantization noise is independent across subspaces, so the total scan-score noise grows as $\sqrt{m}$ while the useful margin between neighbors does not. At $d=200$ ($m=100$ subspaces, wide OOD margins) this is harmless — which is exactly why the int8 path is the right champion default there. At $d=768$ ($m=384$, small in-distribution cosine margins) the noise swamps the signal. The discriminating experiment holds the *index and codes fixed* and changes only the scan arithmetic: int8-saturating scan orders the pool at 0.07; an int16-LUT scan orders it at 0.87 — the pool ceiling — at rank depth 100.

The fix is a policy, not a rewrite. Scan precision is selected by subspace count: $m \leq 128$ keeps the int8 fast-scan (bit-identical to the OOD champion path), $m > 128$ uses int16 LUTs. Two systems consequences follow. First, the survivor cut becomes trustworthy, so the re-rank pool collapses from 32,000 back to a few hundred. Second, the int16 kernel is worth engineering: a 32-wide AVX-512 (*vpermw*) int16 pair-scan — with its dispatch hoisted out of the per-block hot loop, where a per-call environment lookup had been costing $\sim 25\%$ of the scan — is compute-bound at $m=384$ (sorted-vs-scattered block order = $1.00\times$), recall-bit-identical, and $+31\%$ end-to-end. Notably this *reverses* the wider-SIMD negative of §13: at $m=100$ the scan is memory-bound and AVX-512 buys $\leq 7\%$; at $m=384$ it is LUT-compute-bound and the same width is decisive. Precision and width are functions of m , not constants.

Composition. With precision fixed, the OOD coverage levers apply unchanged in-distribution: k -NN-graph augmentation (edges from a base self-search) cuts probes $\sim 3\times$ at matched recall, and a $\text{dpb}=4$ recode ($m=192$: half the LUT work per candidate) trades 0.7 recall points for $\sim 1.3\times$ QPS, net-positive on the frontier.

Effect. Cohere-768 1M, 1 thread: 0.9197@1957 QPS *graph-free* vs. HNSW 0.9122@1840 — a standalone win. At 10M/ $d=768$, HNSW’s strongest regime, HNSW leads by ~ 1.4 – $1.7\times$ (1t) at matched recall once both use *self-built* graphs (Table 7); an earlier 10M lead we measured relied on an HNSW-constructed graph and does not hold standalone (§7).

12 Systems Techniques (Streaming)

We evaluate the streaming track on the msturing-30M-clustered runbook, targeting recall@10 under the 1 h + 8 GB budget. Three levers carry the result. *Batch-parallel insert* parallelizes assignment over each insert batch while preserving order, so the index is bit-identical to serial insertion. *Low-live adaptive p* grows the probe count when the live set is small, so the expected candidate count stays above threshold early in the stream. A *parallelization fix* repairs the batched re-rank driver, which had run chunks serially so that Rayon never engaged. On top of these we add a compliance re-architecture: a first-insert-batch cold start followed by periodic live-set retraining.

Effect. Batch insert: 30M inserts 1818s $\rightarrow$ 406s ($4.5\times$), recall-identical. Adaptive p : worst step 0.737 $\rightarrow$ 0.944. Parallelization fix: $7.6\times$ at 8 threads. Final compliant series (NQ=10000, 640 steps): the budget-eligible frontier on our (shared, 2–3 $\times$ slower than the official VM) machine is recall@10 = 0.8821 at 52.7 min;

Table 3: Streaming (msturing-30M-clustered runbook; NQ=10000, 640 steps) on our shared box under the 8 GB DRAM + 1 h budget. Recall-vs-budget is compute-bound and machine-relative — this box runs $\sim 2\text{--}3\times$ slower than the official 8-vCPU VM, where the same configs project to 0.90–0.92 eligible.

Operating point	recall@10	wall (min)	peak mem
budget-eligible frontier (this box)	0.8821	52.7	7.1–7.2 GB
higher recall (over 1 h here)	0.9010	69.0	7.1–7.2 GB
Enabling levers (§12): batch-parallel insert (30M inserts 1818 $\rightarrow$ 406s, 4.5 $\times$, recall-identical); low-live adaptive p (worst step 0.737 $\rightarrow$ 0.944); a rerank parallelization fix (7.6 $\times$ at 8 threads). Both points stay under the 8 GB cap.			

0.9010 needs 69 min. Recall-vs-budget is compute-bound: eligibility is machine-relative, and the same configs project to 0.90–0.92 eligible on the official 8-vCPU box. Peak memory 7.1–7.2GB < 8GB cap throughout.

13 Negative Results (Evidence for the Localization)

Each of the following was tried and *did not* pay off; we report each with its measured null or loss, since these nulls are what localize the bottleneck to routing and coverage.

- **Query-aware routing** (train cells on the query distribution): *worse* than base-trained (0.76 vs 0.89 @ matched p) — undertrained + the cells where queries are dense don’t distribute base points well.
- **ADC routing** (4-bit-scan the centroids): recall-neutral but *no speed win* — the exact VNNI router is already at the compute floor.
- **Anisotropic/OPQ-rotated coarse codes for rank preservation**: do *not* make coarse codes rank-preserving; reorder depth is set by code bit-rate, and rotation/ η move it < 0.01. Coverage-via-coarsening collapses (a scann-coverage-matched dpb=50 config yields recall 0.09).
- **RaBitQ**: no ranking edge over PQ at matched bytes.
- **Wider SIMD** (AVX-512 vpermw / 64-wide): *on OOD* $d=200$ the scan is memory-bound, not shuffle-bound; $\leq +7\%$, sometimes negative (downclock/cross-lane). The *same* lever is decisive at $m=384$ (§11) — the negative is regime-scoped, not universal.
- **Finer cells at matched probe**: lose recall (half the probe fraction).
- **PCA/LeanVec dim reduction on OOD**: near-isotropic embeddings; dead. The same holds for *raw-basis* routing-dimension truncation in-distribution (Cohere-768: scoring a 384-dim prefix costs 1.6 recall points) — prefix truncation needs a variance-ordering rotation to be principled, in either regime.
- **Residual-encoded PQ under inner product** (baseline fairness): FAISS IVFPQ defaults by_residual=True, which is L_2 -specific; under IP it caps text2image recall at 0.11, *decreasing* with probes. All FAISS-IVFPQ baselines here use by_residual=False plus exact refine — reporting the default would be a strawman.
- **Int8-prefilter + int16-rescore hybrid scan**: dominated at $m=384$ — the int8 pass’s ordering is too noisy for any affordable keep rate (top-2000 of 8000 retains the true neighbor only 53% of the time), so the prefilter cannot shrink the int16 work without recall loss.

The OOD negatives cluster on the scan/code side and all fail there; the §5–9 levers cluster on routing/coverage and all succeed. In-distribution the pattern inverts: coverage was never the binding constraint, scan precision was (§11). Both localizations rest on the same isolation method.

14 Evaluation Methodology

We report *same-hardware ratios*, never cross-machine QPS. All comparisons run on one physical machine; each competitor is measured in the same session, seconds apart, at its own best operating point.

Table 4: OOD text2image, recall@10-matched QPS ratio (ScaNN/ours; < 1 = our engine faster). Same-hardware, tight-pairwise, opponent at official config. At 100M no ScaNN index builds within the gate (note), so the comparison is on build-eligibility.

Scale	1-thread	8-thread	recall@10 gate
1M	0.755	0.746	0.905
10M	0.827	0.758	0.901
100M	ScaNN build-ineligible (> 2 h gate; see note)		

At 100M no ScaNN configuration builds within the 2h gate (40k-leaf killed at 121 min; a coarser 20k-leaf also exceeds 2h — the $O(100M)$ AH-encode/bf16 prep dominates regardless of leaf count), so no matched-recall ratio exists. Ours builds a 0.903-recall index in 86 min and serves ~ 6.7 k QPS at recall 0.90 (16 threads) — under the 2h gate it is the only method with a 100M index.

Protocol. For a target recall ($\text{recall@10} \geq 0.90$), we (i) fix each system’s operating point by an independent sweep, requiring both to clear the recall gate against the *same* exact float ground truth; (ii) alternate the two systems round to round (order flips) so neither owns a load window; (iii) take best-of- k within a round and report the *median of per-round ratios*; (iv) recompute recall from raw ids, never trusting a cached number. On a shared machine, load drift between two builds minutes apart exceeds the effects being measured, so only interleaved same-window ratios and load-independent recall are trusted.

Warm-resident semantics; pristine confirmation. Both systems measured warm and resident (leaderboard serving); an early cold-spawn asymmetry is corrected by scoring each engine’s warm block. Headline numbers pause all other first-party jobs (SIGSTOP) and confirm across ≥ 20 rounds.

The baseline at its best. At 10M, where an *official* big-ANN recipe exists, ScaNN uses it verbatim (40k leaves, anisotropic AH/LUT16, residual quant, bfloat16 re-order, top-level partitioner) with its own leaves-to-search / re-order sweep including corner configs. At scales with no official recipe (1M, 100M) we instead give ScaNN its *speed-optimal* configuration, found by sweeping the leaf count: e.g. at 1M, ScaNN peaks at 1200 leaves (6.2k QPS@0.90), with both finer and coarser partitions slower — coarser because routing over more leaves outweighs the finer scan, mirroring our own coarse-cell result (§8). This matters: a convention-chosen leaf count can silently under- or over-partition the baseline. A win depending on a handicapped baseline is quarantined; we describe one (an under-leaved cached 10M index) and its correction.

Build-time eligibility. The benchmark caps index build at **12 hours** on the 8-vCPU evaluation machine, so a baseline (or our own index) that cannot build in time is not merely slow — it is ineligible. This is a real constraint at 100M: a $\sqrt{n}$ -scaled 126k-leaf ScaNN would take ≈ 32 h at 8 vCPU (its partitioner k -means over the training sample), so the eligible ScaNN baseline is a coarser config that builds in time. Our indices build comfortably inside the limit and *faster than ScaNN at scale*: on 16 cores, 1M in 131s, 10M in 27min, 100M in 86min (≈ 3 h at 8 vCPU), versus ScaNN’s ≈ 1.5 h (10M) and ≈ 2 h (100M, eligible config). We therefore report build time alongside QPS (Table 4); the hierarchical- k -means+PQ build is cheaper than ScaNN’s partitioner- k -means+AH training, and the gap widens with n .

What we do not claim. Same-hardware ratios on one (shared) machine, not a standardized-VM submission; ScaNN is the *measured* opponent — gaps to binary-only leaderboard systems are *inference* from the public leaderboard.

15 Results

16 Related Work

Quantization-side IVF. ScaNN introduced anisotropic vector quantization (score-aware loss) and remains the measured opponent here at its official configurations; FAISS provides the reference IVF-PQ/fast-scan

Table 5: Ablation ladder at 10M (ScaNN/ours ratio; each row adds one lever).

Configuration	ratio
γ only	1.24
+ graph coverage	1.01
+ tree-EM	0.92
+ coarse-cell geometry (knee)	0.85
+ pristine confirmation (20 rounds)	0.83

Table 6: Multi-system baselines, OOD text2image-10M, 1 thread, same box, same exact ground truth. QPS at the listed recall@10 (each system swept to its knee; FAISS-IVFPQ uses `by_residual=False` + exact refine, its fair IP configuration). RoarGraph [5] is the published OOD-specific graph method, built with its paper hyperparameters ($M_{sq}=100$, $M=35$, $L=500$, 2M sampled train queries, GT recall@100 0.96).

System	recall@10	QPS (1t)	note
ours	0.905	7657 (1M) / 2804 (10M)	IVF-PQ cascade
ScaNN (official 40k)	0.901	(ratio 0.827 vs ours)	IVF-AH
RoarGraph	0.903	2250 [‡]	OOD bipartite graph
HNSW (faiss, M32)	0.905	328	needs $ef \geq 160$ on OOD
IVFPQ-FastScan+refine	cannot reach 0.90 (caps 0.73)		code-side fails OOD

[‡]5-round same-window pairwise median vs ours at matched recall: ours/RoarGraph = 1.07 $\times$.

implementations (and §13 documents an IP-fairness pitfall in its residual default). RaBitQ represents the recent binary-ish code family; at matched bytes it showed no ranking edge in our OOD setting. Our contribution on this side is not a new code but the observation that *scan arithmetic precision must scale with code length* (§11). **Graph indices.** HNSW is the standard in-distribution baseline and is measured here at both regimes: dominant-until-matched in-distribution, and $>4\times$ behind on OOD at matched recall. RoarGraph (VLDB’24) is the OOD-specific graph method — a query-aware bipartite projection built from sampled training queries; we build it with its published hyperparameters and measure it on the same box (Table 6). DiskANN/Vamana target the disk regime we do not evaluate. **OOD-specific routing.** Query-aware cell *training* fails in our hands (§13); RoarGraph solves coverage with a query-aware graph; our γ recalibration (§6) instead keeps base-trained cells and corrects the *scoring* of them, at zero query-time cost, composing with graph coverage rather than replacing the IVF skeleton. **Streaming.** The big-ANN streaming track constraints (1h wall, 8GB) shape §12; our levers there are systems-level (batch-parallel insert, adaptive probes) rather than index-structural.

References

- [1] R. Guo et al. Accelerating large-scale inference with anisotropic vector quantization. *ICML*, 2020.
- [2] P. Sun, R. Guo, S. Kumar. SOAR: improved indexing for approximate nearest neighbor search. *NeurIPS*, 2023.
- [3] Yu. Malkov, D. Yashunin. Efficient and robust approximate nearest neighbor search using hierarchical navigable small world graphs. *TPAMI*, 2020.
- [4] M. Douze et al. The Faiss library. *arXiv:2401.08281*, 2024.
- [5] M. Chen et al. RoarGraph: a projected bipartite graph for efficient cross-modal approximate nearest neighbor search. *PVLDB*, 2024.
- [6] S. Jayaram Subramanya et al. DiskANN: fast accurate billion-point nearest neighbor search on a single node. *NeurIPS*, 2019.
- [7] W. Dong, M. Charikar, K. Li. Efficient k -nearest neighbor graph construction for generic similarity measures. *WWW*, 2011.

Table 7: In-distribution Cohere-768 (cosine, unit-norm), 1 thread, same box, same exact ground truth. On the shared machine only interleaved same-window ratios are trusted (§14); at 10M we report the 5-round tight-pairwise median ratio and the cleanest-window absolute pair. Recall is load-independent.

Scale	System	recall@10	QPS (1t)
1M	ours (int16, graph-free)	0.9197	1957
1M	HNSW (ef40)	0.9122	1840
1M	ScaNN (2k leaves, reorder 600)	0.9147	1306
1M	IVFPQ+refine (nprobe 256)	0.9014	877
10M	ours (<i>self-built</i> graph, hops=2/3)	0.9707 / 0.9832	1234 / 781
10M	HNSW (ef96 / ef160)	0.9677 / 0.9803	772 / 467
10M	ours (<i>self-built</i> graph, hops=3, $p=128$; top)	0.9904	483
10M	HNSW (ef240; its ceiling)	0.9859	308
10M	ours (<i>routed</i> self-search graph, hops=3)	0.9619	751
10M	ScaNN (8k leaves, deep reorder)	0.9461	~330

At **1M** the engine wins *graph-free* and standalone. At **10M**/ $d=768$ (HNSW’s strongest regime) the engine beats HNSW at every operating point using its *self-built* NN-descent graph (clean-window sequential runs, same cores): $\sim 1.27\times$ at recall 0.946, $\sim 1.6\times$ at 0.97, $\sim 1.67\times$ at 0.983, $\sim 1.9\times$ at HNSW’s 0.9859 ceiling (1 thread; $1.2\text{--}1.6\times$ at 8 threads), higher recall throughout, and it reaches recall 0.99 where HNSW plateaus (≈ 0.986). The graph *need not be near-exact*: degrading its overlap@16 from 0.93 to 0.65 costs only $\sim 0.4\text{pt}$ recall at matched probes, and genbo still beats HNSW on *both* recall and QPS with a 0.65-overlap graph (at a modestly higher probe count); even a 0.28-overlap graph recovers most of the coverage benefit. QPS is nearly flat across graph quality — the graph buys coverage/recall, not speed. And the graph is *self-built*: the engine’s internal NN-descent pass (§7) reaches 0.904 overlap in minutes at 10M, and the resulting self-built-graph rows are query-time equal-or-better than the external near-exact graph. (The *routed* self-search variant, last row, is the one that fails — it is coverage-capped at $d=768$ and loses to HNSW; NN-descent does not route, which is the point.) ScaNN trails throughout.

- [8] J. Gao, C. Long. RaBitQ: quantizing high-dimensional vectors with a theoretical error bound for approximate nearest neighbor search. *SIGMOD*, 2024.
- [9] H. Simhadri et al. Results of the big-ANN: NeurIPS’23 competition. *arXiv:2409.17424*, 2024.
- [10] F. André, A.-M. Kermarrec, N. Le Scouarnec. Quicker ADC: unlocking the hidden potential of product quantization with SIMD. *TPAMI*, 2021.

17 Limitations

In-distribution at scale: a graph-build premium (now internal). At Cohere 1M/ $d=768$ the engine wins *graph-free*. At 10M/ $d=768$ (HNSW’s strongest regime) its query path beats HNSW across the recall range ($\sim 1.2\text{--}1.9\times$, higher recall throughout, and it reaches 0.99 where HNSW plateaus at ≈ 0.986 ; Table 7) — but only with an IP- k -NN graph. That graph *need not* be near-exact: we degraded the 0.93-overlap graph and find the win survives down to ~ 0.65 overlap (genbo still beats HNSW on both recall and QPS, at a modestly higher probe count), with the recall benefit saturating quickly — even a 0.28-overlap graph recovers most of it. Crucially, the graph no longer requires an external builder: genbo’s *routed* self-search is coverage-capped at this dimension (it reaches only ~ 0.53 exact-overlap@16, and even probing 6% of cells lifts this to just ~ 0.63 — the true neighbours are scattered across too many Voronoi cells for any practical probe count to cover), but its internal *NN-descent* pass (§7) does not route at all and self-builds a 0.904-overlap graph in minutes at 10M — query-time performance equal-or-better than the multi-hour external near-exact graph. What remains is honest accounting of the *build-time premium*: the descent pass adds minutes to the build (random init 24 min for 0.864 overlap; seeded 8 min for 0.904, at 16 threads), a cost graph methods like DiskANN/NSG already assume as their core build step. The reason genbo needs the graph at all is a *candidate-efficiency* gap we profiled: graph-free, to reach recall 0.96 it scans and exactly-rescores ~ 1600 candidates/query (memory-latency-bound $\sim 0.6\mu\text{s}/\text{row}$, already prefetched) where HNSW’s greedy traversal touches fewer — the graph supplies exactly the coverage the routing misses. This need does *not* persist at higher dimension: with multi-hop graph repair (§7) at 35M/ $d=1024$ (Wikipedia Cohere-v3, unit-norm,

Table 8: In-distribution Wikipedia Cohere-v3 (35M, $d=1024$, unit-norm, strongly anisotropic $\|\mu\|=0.48$), 1 thread, same box. Ground truth is recomputed over the queries-excluded base and HNSW is rebuilt on that same base (int8 scalar-quantized storage), so both systems search identical data (§17). Recall is load-independent.

Recall band	System	recall@10	QPS (1t)
~0.966	ours (dpb4 + graph hops=2, $p=48$)	0.9664	453
~0.966	HNSW (ef160)	0.9647	321
~0.982	ours (dpb4 + graph hops=2, $p=96$)	0.9817	306
~0.982	HNSW (ef640, its ceiling)	0.9818	95
top	ours (hops=3, $p=96$)	0.9827	304
top	HNSW (max, ef $\gg$ 640 saturates)	≈ 0.982	—

The graph here is built by ours’ OWN self-search (**selfknn**, §7), not borrowed from a competitor — the comparison is self-contained. At matched recall ours is $\sim 1.4\times$ (recall 0.966) and $\sim 3\times$ (0.982) faster than HNSW single-thread, and reaches recall 0.983 that HNSW cannot; the win holds at 8 threads ($1.16\text{--}2.8\times$). HNSW is faster only below recall ≈ 0.95 . (An HNSW-built graph lifts ours by ~ 0.4 pt more, but we report the self-built result.)

strongly anisotropic $\|\mu\|=0.48$) genbo *beats* a memory-matched HNSW across the practical high-recall range using its own self-built graph ($1.4\text{--}3\times$ at recall ≥ 0.96 , reaching recall HNSW cannot; Table 8), HNSW faster only below ≈ 0.95 — there HNSW is weaker (it plateaus near 0.982) and genbo’s own graph nearly matches a competitor-built one. **Benchmark hygiene (a cautionary note).** An earlier version of this dataset appeared to cap our recall at 0.62 while HNSW reached 0.97; the cause was a ground-truth construction pitfall, not an algorithmic boundary. When the query set is carved from the base and the exact ground truth is computed over the base *including* those rows, an index that correctly *excludes* the queries faces an artificial recall ceiling—here 38% of each top-10 were query self-matches or near-duplicate query rows, capping recall at 0.62—while an index that inadvertently *retains* the queries appears to win. All numbers here use ground truth recomputed over the queries-excluded base, with every method searching that same base. We flag this because the failure is silent: both indices run without error and the ceiling looks like a model deficiency. **Other scope:** single (shared) machine, not a standardized-VM submission (the gating experiment for a leaderboard claim); γ and scan-precision are fit/selected per dataset; the coarse-cell corollary is shown on IP/OOD data; the gap to binary-only leaderboard systems is inference from public numbers; filter/sparse tracks are untouched.